Christian Bourlier

STATUS: ACTIVE // ROLE: AI_SYSTEMS_ARCHITECT // CLASS: APPLIED_AI_&_AUTOMATION

github.com/rezzedai · christian@bourlier.ai · bourlier.ai

[BENCHMARKS]

3	39	100%	0	6mo→6wk
PROD AI SYS	MCP TOOLS (9 DOMAINS)	AUDIT COVERAGE	UNCONTAINED FAILURES	MODELING LATENCY REDUCTION

> *"Christian brings both technical excellence and great team spirit...*
> *He picks up complex topics fast and always delivers on goals."*
>
> — Brianna Mersey, VP Data

HOW I THINK

Systems First
Every problem is a system. Map inputs, outputs, feedback loops, constraints, and failure modes before code. Build the smallest architecture that satisfies them.

Typed Work > Vanity Metrics
Orchestration differentiates WORK from CONTROL signals. Measurement clarity prevents metric inflation.

Constrained Autonomy   [GOVERNANCE]
Agents operate within identity boundaries, context budgets, rate limits, and pre-execution policy gates. Safety, evaluation, and dispatch layers enforce bounded execution.

THE APPROACH

- 01 **DIAGNOSE** Map the system
- 02 **BOUND** Define constraints
- 03 **ORCHESTRATE** Wire the agents
- 04 **INSTRUMENT** Measure everything
- 05 **DEPLOY** Ship to production
- 06 **ITERATE** Eval-driven loops

INTERESTED IN

Applied AI systems where autonomy is governed by explicit boundaries and measurable outcomes — orchestration, deployment under real-world constraints, and evaluation-driven iteration.

CERTIFICATIONS
- GCP Professional Data Engineer
- GCP Professional Cloud Architect

PRODUCTION SYSTEMS

Rezzed.ai
`bounded-multi-agent-orchestration`
Claude Code. Role-specialized agents in Planner □ Worker topology. Typed task lifecycles (WORK vs CONTROL), identity-scoped boundaries, context monitoring, static token ceilings, pre-tool-use policy gating. Vendor-neutral backends.

CacheBash
`mcp-relay + persistent-task-state`
Cloud Run, transactional claiming. TTL-based relay, dead-letter containment, identity isolation, audit logging at tool boundary. Separates productive work from control traffic.

OptiMeasure
`privacy-first-ai-measurement`
Vertex AI cookieless attribution. Modeling latency: 6 months □ 6 weeks. Clean-room compatible.

Bourlier.ai
`constraint-aware-deployment`
Production testbed. Gemini + rate limiting, validation, audit logging, structured output handling.

`// instrumentation.underway`

- Completion semantics (SUCCESS / FAILED / SKIPPED)
- Retry classification + error taxonomy
- Token, cost, latency rollups per task class
- Prompt/output provenance hashing

```yaml
# system_manifest.yaml
candidate: Christian Bourlier
role: "AI Systems Architect"
core_logic: "Bounded Autonomy"
primary_stack:
  - Claude Code
  - Vertex AI
  - Cloud Run + MCP
safety: "Constrained Execution"
status: "production_ready"
```