

Christian Bourlier

STATUS: ACTIVE // ROLE: AI_SYSTEMS_ARCHITECT // CLASS: APPLIED_AI_&_AUTOMATION
github.com/rezzedai · christian@bourlier.ai · bourlier.ai

[BENCHMARKS]

3	39	100%	0	6mo→6wk
PROD AI SYS	MCP TOOLS (9 DOMAINS)	AUDIT COVERAGE	UNCONTAINED FAILURES	MODELING LATENCY REDUCTION

"Christian brings both technical excellence and great team spirit... He picks up complex topics fast and always delivers on goals."

— Brianna Mersey, VP Data

HOW I THINK

Systems First

Every problem is a system. Map inputs, outputs, feedback loops, constraints, and failure modes before code. Build the smallest architecture that satisfies them.

Typed Work > Vanity Metrics

Orchestration differentiates WORK from CONTROL signals. Measurement clarity prevents metric inflation.

Constrained Autonomy

[GOVERNANCE]

Agents operate within identity boundaries, context budgets, rate limits, and pre-execution policy gates. Safety, evaluation, and dispatch layers enforce bounded execution.

THE APPROACH

- ① **DIAGNOSE** Map the system
- ② **BOUND** Define constraints
- ③ **ORCHESTRATE** Wire the agents
- ④ **INSTRUMENT** Measure everything
- ⑤ **DEPLOY** Ship to production
- ⑥ **ITERATE** Eval-driven loops

INTERESTED IN

Applied AI systems where autonomy is governed by explicit boundaries and measurable outcomes – orchestration, deployment under real-world constraints, and evaluation-driven iteration.

CERTIFICATIONS

- GCP Professional Data Engineer
- GCP Professional Cloud Architect

PRODUCTION SYSTEMS

Rezzed.ai

bounded-multi-agent-orchestration
Claude Code. Role-specialized agents in Planner □ Worker topology. Typed task lifecycles (WORK vs CONTROL), identity-scoped boundaries, context monitoring, static token ceilings, pre-tool-use policy gating. Vendor-neutral backends.

CacheBash

mcp-relay + persistent-task-state
Cloud Run, transactional claiming. TTL-based relay, dead-letter containment, identity isolation, audit logging at tool boundary. Separates productive work from control traffic.

OptiMeasure

privacy-first-ai-measurement
Vertex AI cookieless attribution. Modeling latency: 6 months □ 6 weeks. Clean-room compatible.

Bourlier.ai

constraint-aware-deployment
Production testbed. Gemini + rate limiting, validation, audit logging, structured output handling.

// instrumentation.underway

- Completion semantics (SUCCESS / FAILED / SKIPPED)
- Retry classification + error taxonomy
- Token, cost, latency rollups per task class
- Prompt/output provenance hashing

```
# system_manifest.yaml
candidate: Christian Bourlier
role: "AI Systems Architect"
core_logic: "Bounded Autonomy"
primary_stack:
  - Claude Code
  - Vertex AI
  - Cloud Run + MCP
safety: "Constrained Execution"
status: "production_ready"
```